



Building event-driven architectures on AWS

Peterson Larentis

Sr. Specialist Solutions Architect, Serverless - LATAM
AWS

Who AMI?



/peterson-larentis



Agenda

Challenges with distributed systems

Journey to event-driven architectures

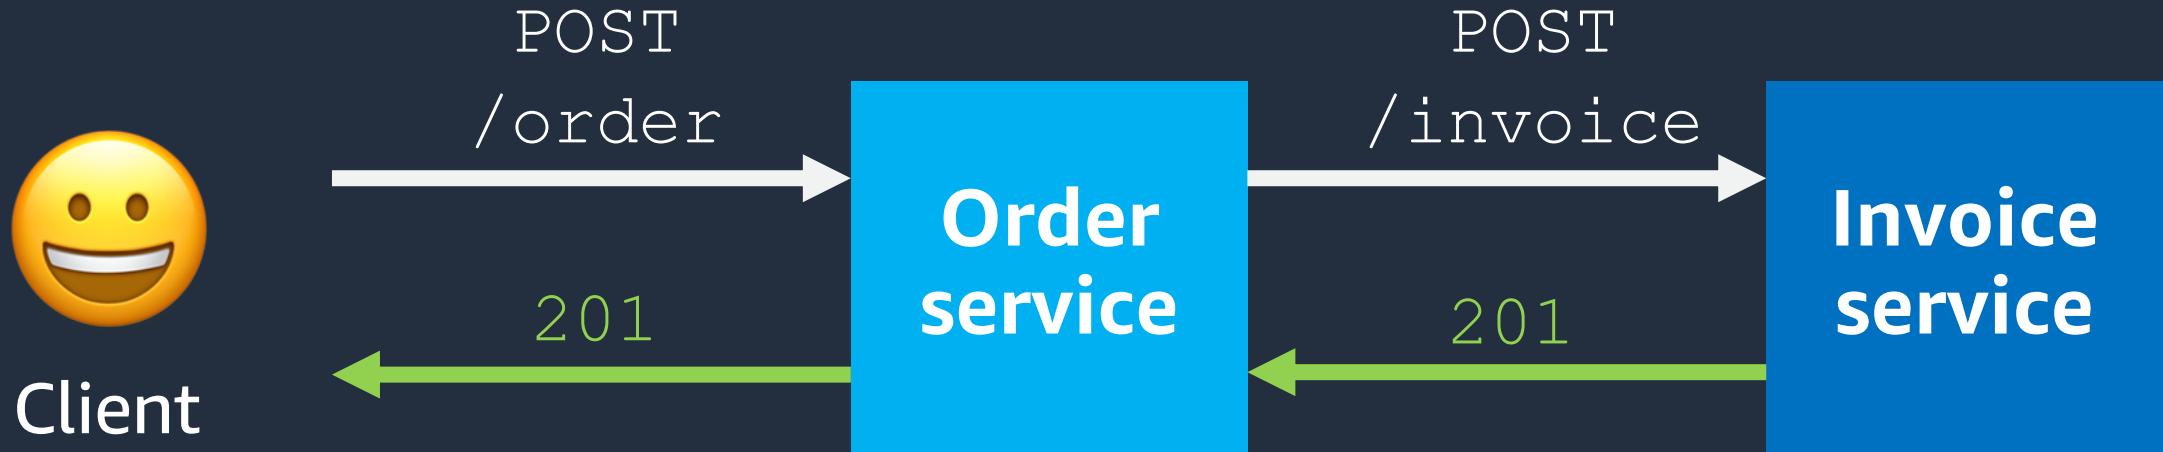
Design considerations

Next steps

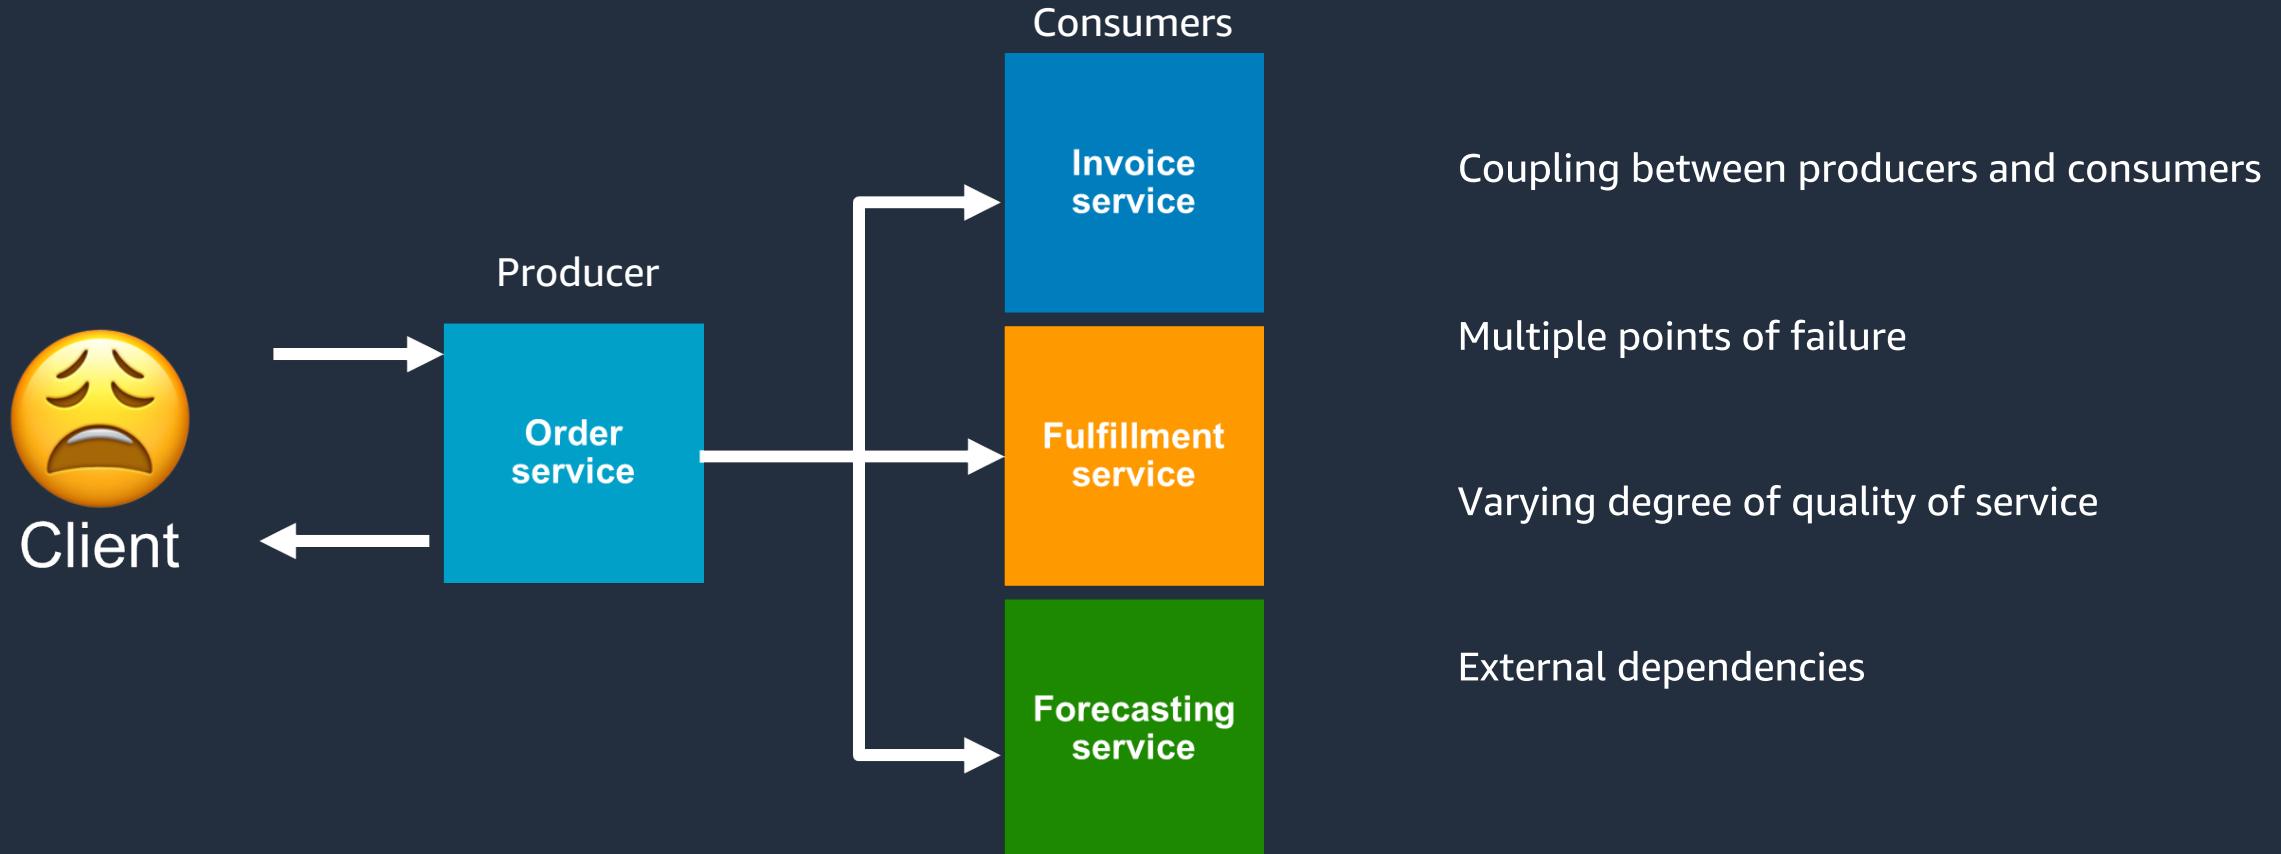
Challenges with distributed systems



Microservices start simple



Synchronous API challenges over time



Coupling – Integration's magic word



Coupling is a measure of independent variability between connected systems.

Decoupling has a cost, both at design and run-time.

Coupling isn't binary.

Coupling isn't one-dimensional.

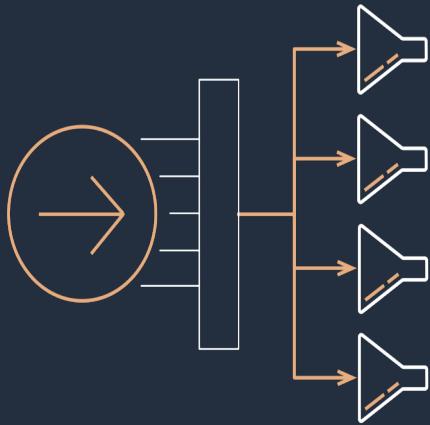
Source: EnterpriseIntegrationPatterns.com

The many facets of coupling

Temporal dependency:	sync, async
Location dependency:	IP addresses, DNS
Data format dependency:	Binary, XML, JSON, ProtoBuf, Avro
Semantic dependency:	Name, Middlename, ZIP
Data type dependency:	int16, int32, string, UTF-8, null, empty
Interaction style dependency:	messaging, RPC, query-style (GraphQL)
Conversation dependency:	pagination, caching, retries
Technology dependency:	Java vs. C++

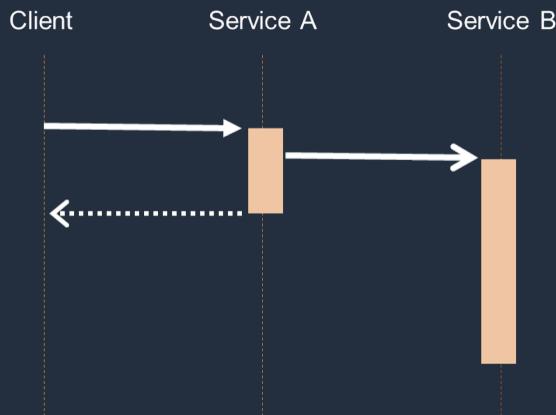


Event-driven architectures drive reliability and scalability



Event Routers

Abstract producers and consumers from each other



Asynchronous Events

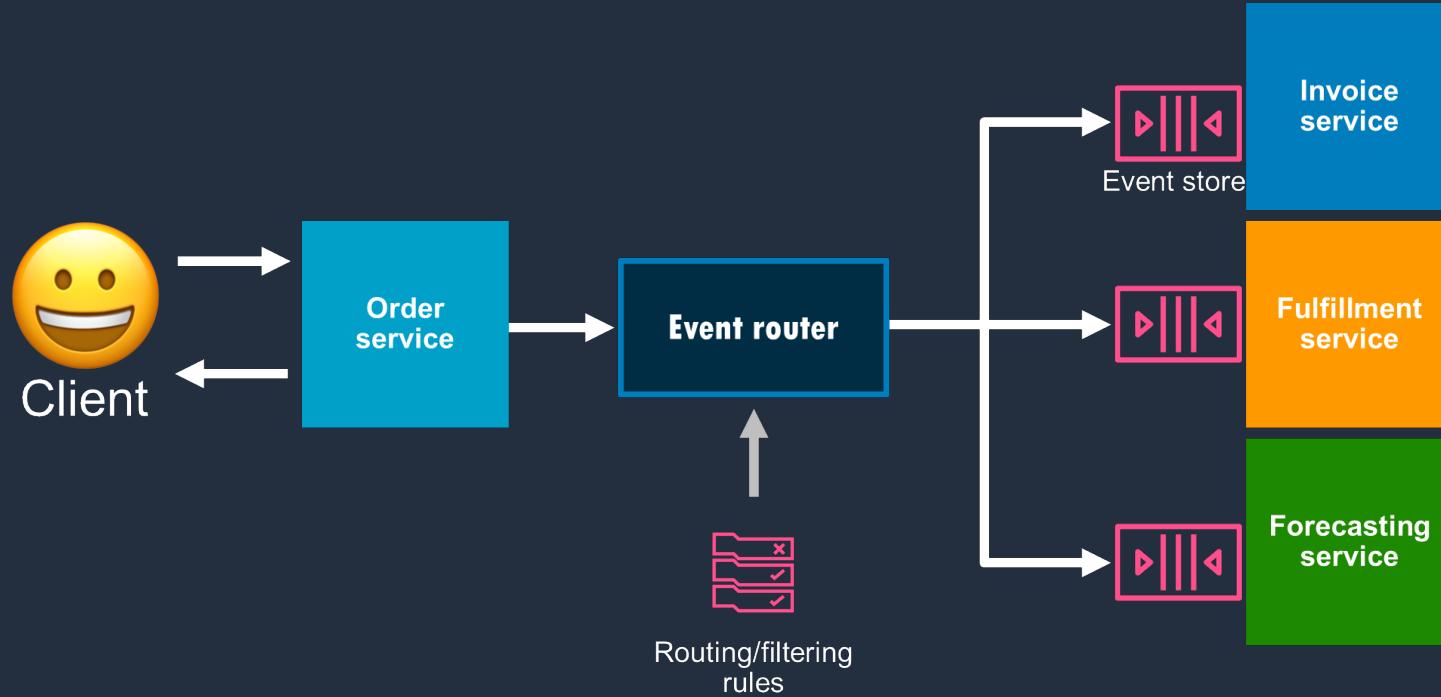
Improve responsiveness and reduce dependencies



Event Stores

Buffer messages until services are available to process

Reliable, resilient, and independently scalable



If your application is cloud-native, or large-scale, or distributed, and doesn't include a messaging or event component, that's probably a bug.

Tim Bray
General-purpose Internet-software geek

Journey to event-driven architectures



Journey to event-driven architectures

STEP 1

Start with the
domain





event

[i-'vent] noun

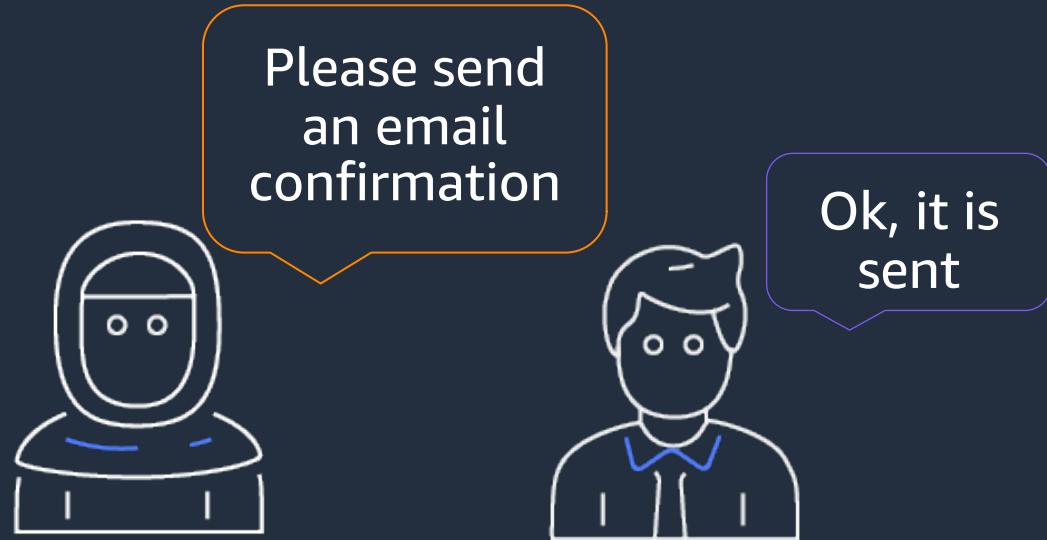
**A signal that a
system's state has
changed.**

Properties of events

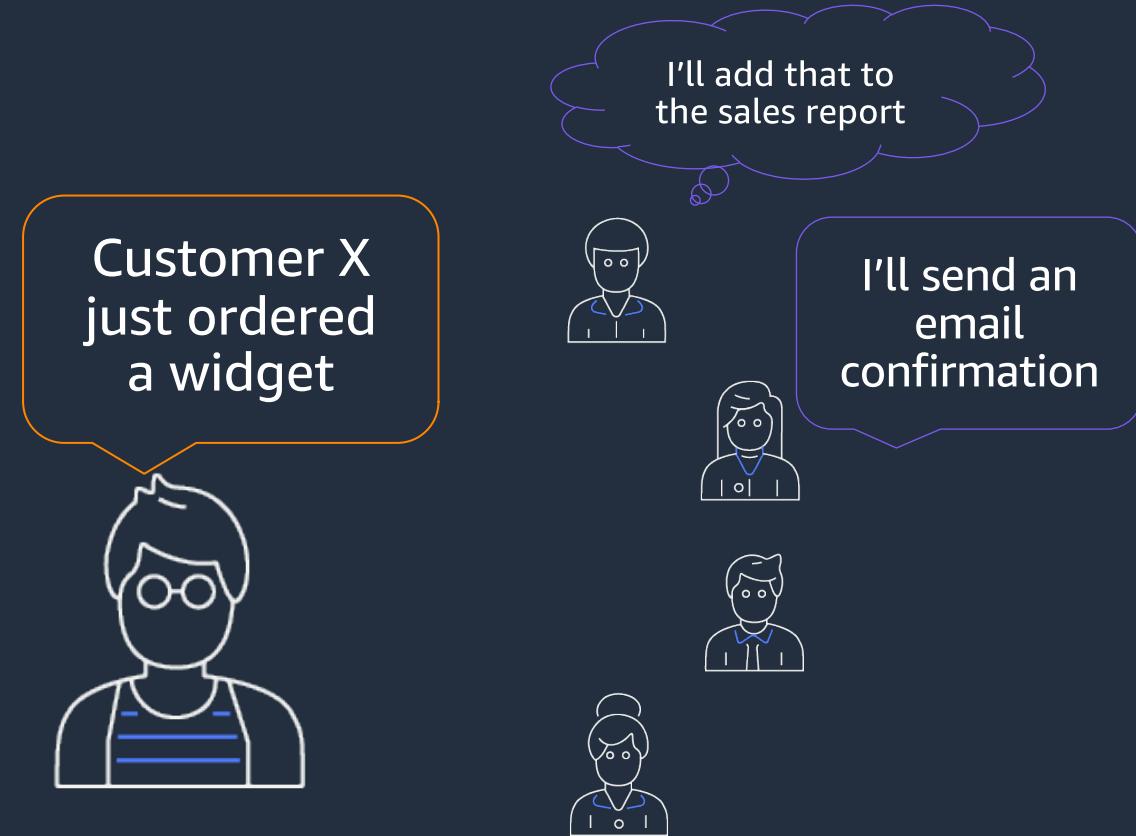
```
{  
  "source": "com.orders",  
  "detail-type": "OrderCreated",  
  "detail": {  
    "metadata": {  
      "idempotency-key": "c1b95b88"  
    },  
    "data": {  
      "order-id": "1073459984"  
    }  
  }  
}
```

- Events are signals that a system's state has changed
- Events occur in the past (e.g. OrderCreated)
- Events cannot be changed (immutable)
- Decrease semantic coupling by restricting information to key data

Events are observable, not directed



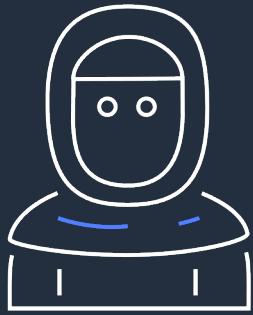
Directed commands



Observable events

Event-driven architectures start with **business events**

OrderCreated



OrderPicked



OrderShipped



ReturnRequested



ReturnReceived

Map events to **business domains** using **attributes**, not entities

OrderCreated



Retail

OrderPicked

OrderShipped



Fulfillment

ReturnRequested

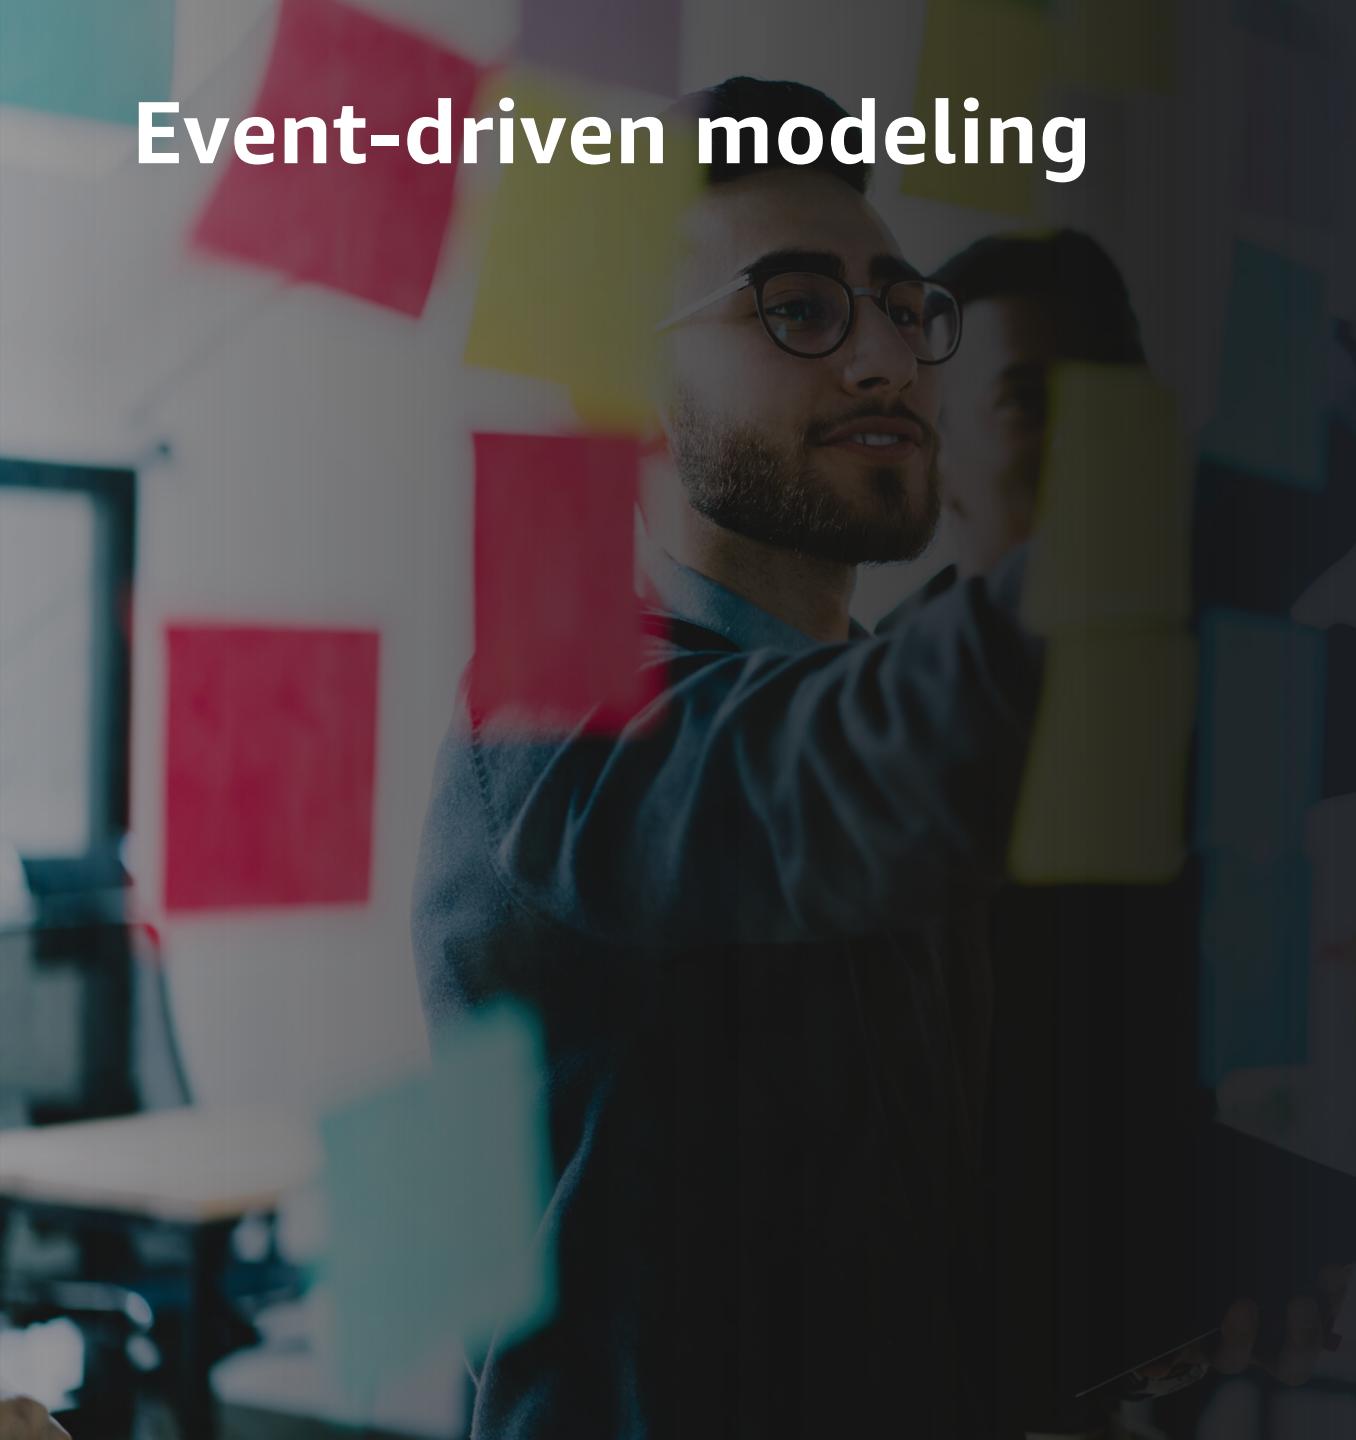


ReturnReceived



Support

Event-driven modeling

- 
- A photograph of a man with dark hair and a beard, wearing glasses and a blue button-down shirt. He is looking towards the right side of the frame with a thoughtful expression. Behind him is a wall covered with numerous colorful sticky notes in various colors like red, yellow, green, and blue, suggesting a collaborative workspace or planning session.
1. Identify business events, processes, actors, etc.
 2. Cluster-related concepts
 3. Define bounded contexts and sub-domains

Journey to event-driven architectures

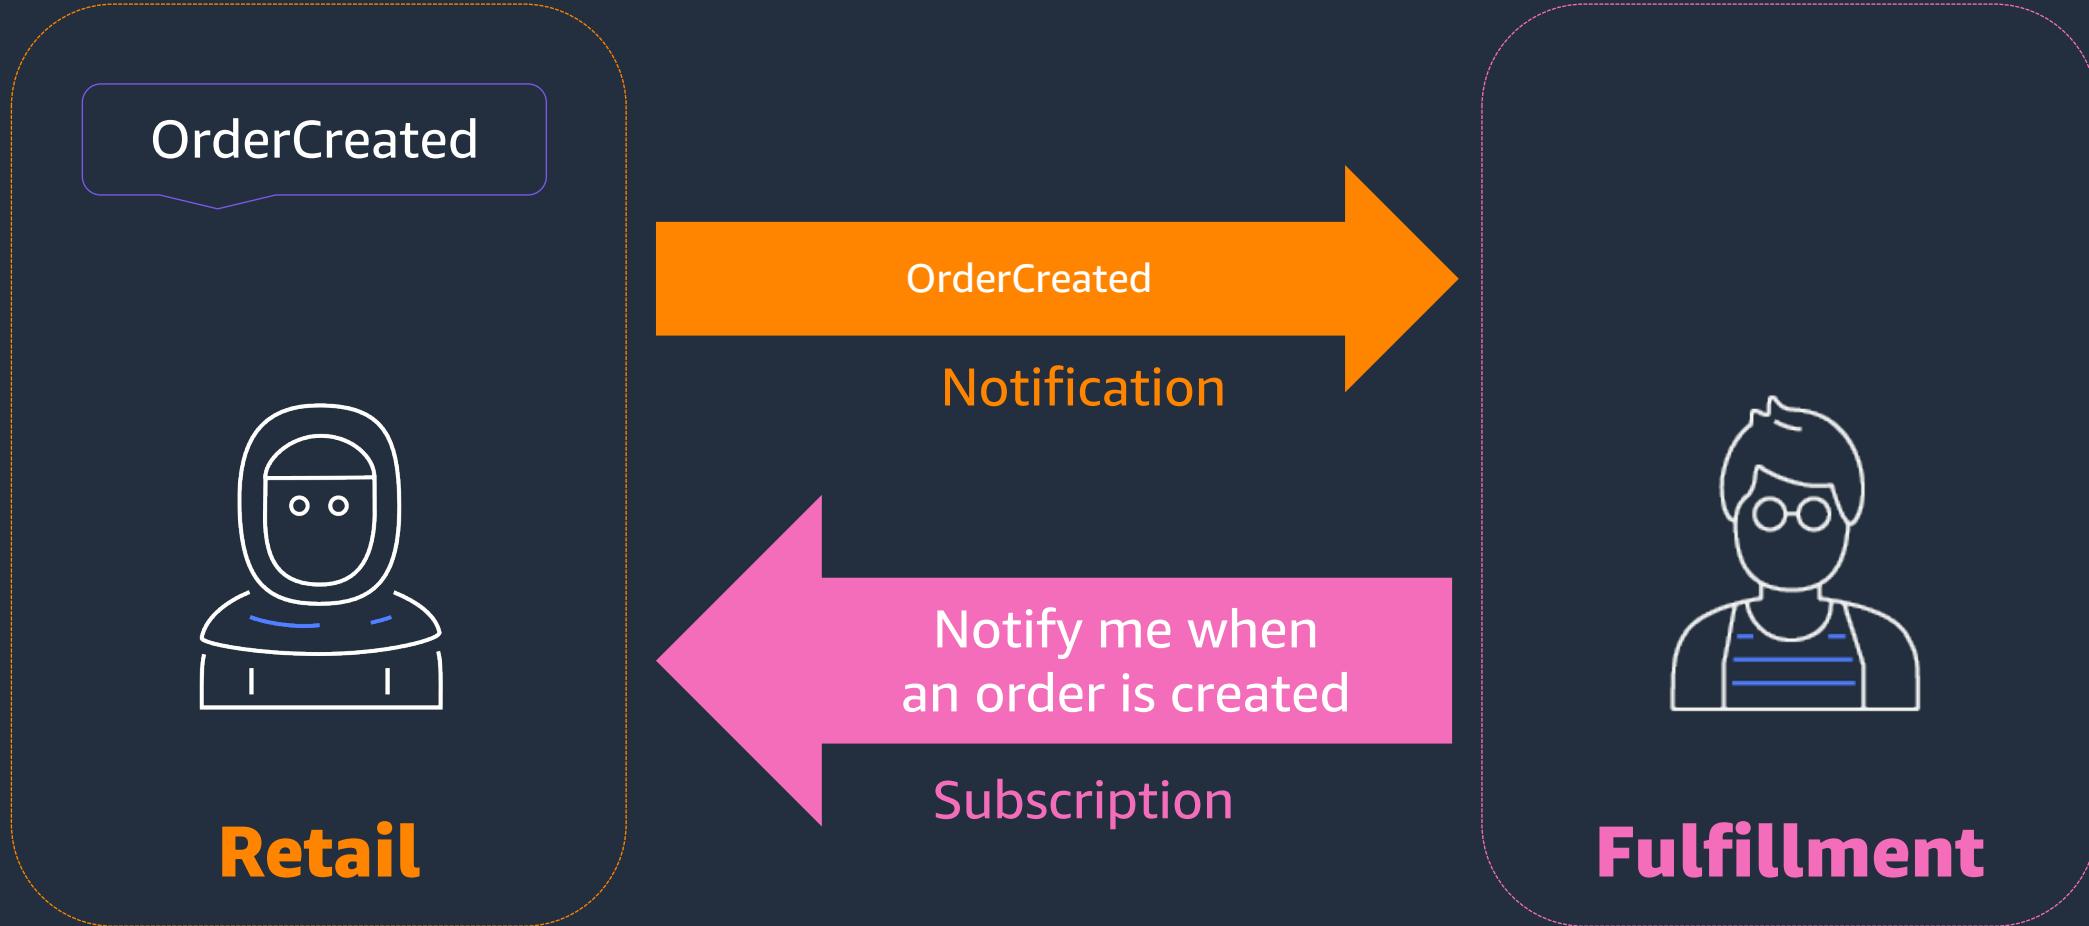
STEP 1

Start with the
domain



STEP 2
Coordinate events

Choreograph events *between domains* using subscriptions



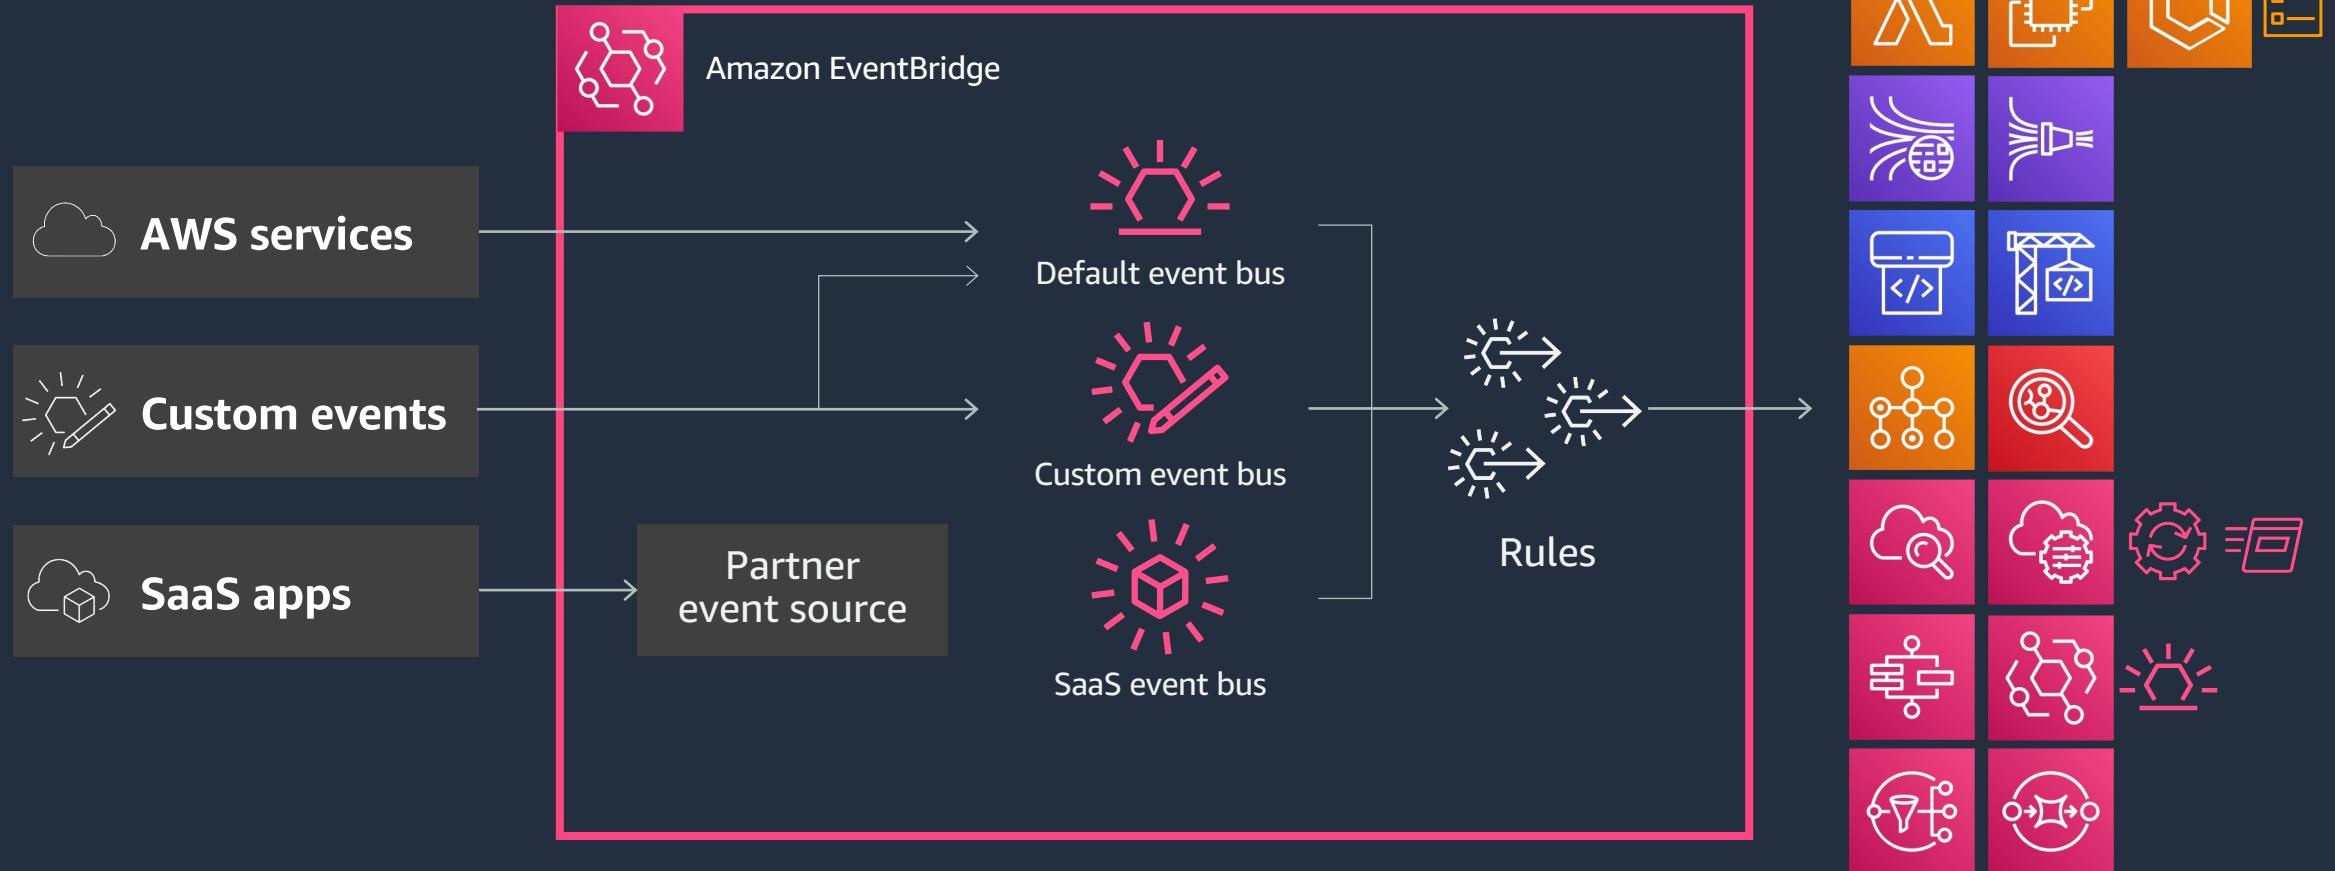


Amazon EventBridge

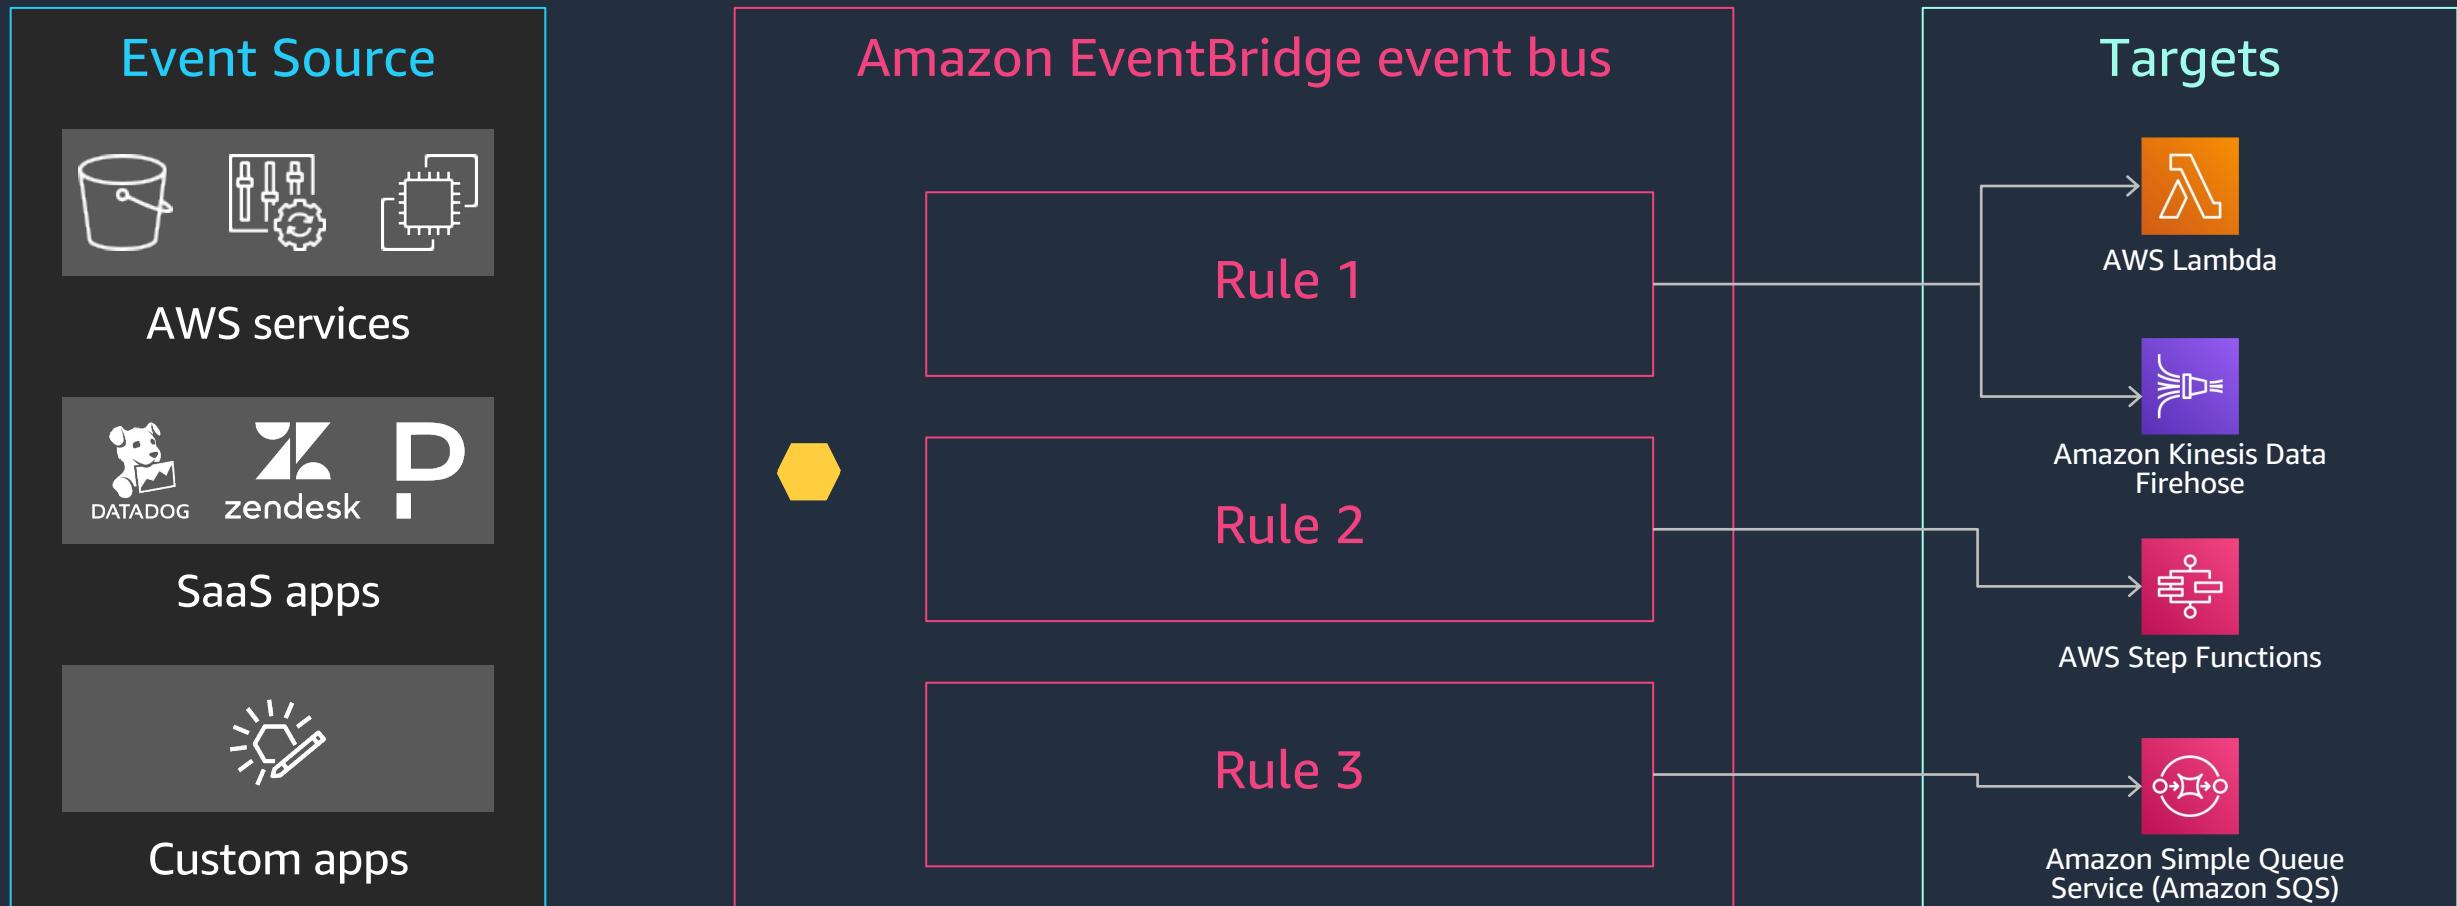
A serverless event bus service for AWS services, your own applications, and SaaS providers.

- Serverless. Pay only for the events you process
- Simplified scaling - avoids increasing costs to sustain and manage resources
- No upfront investments, ongoing licensing, or maintenance costs
- No specialist knowledge needed

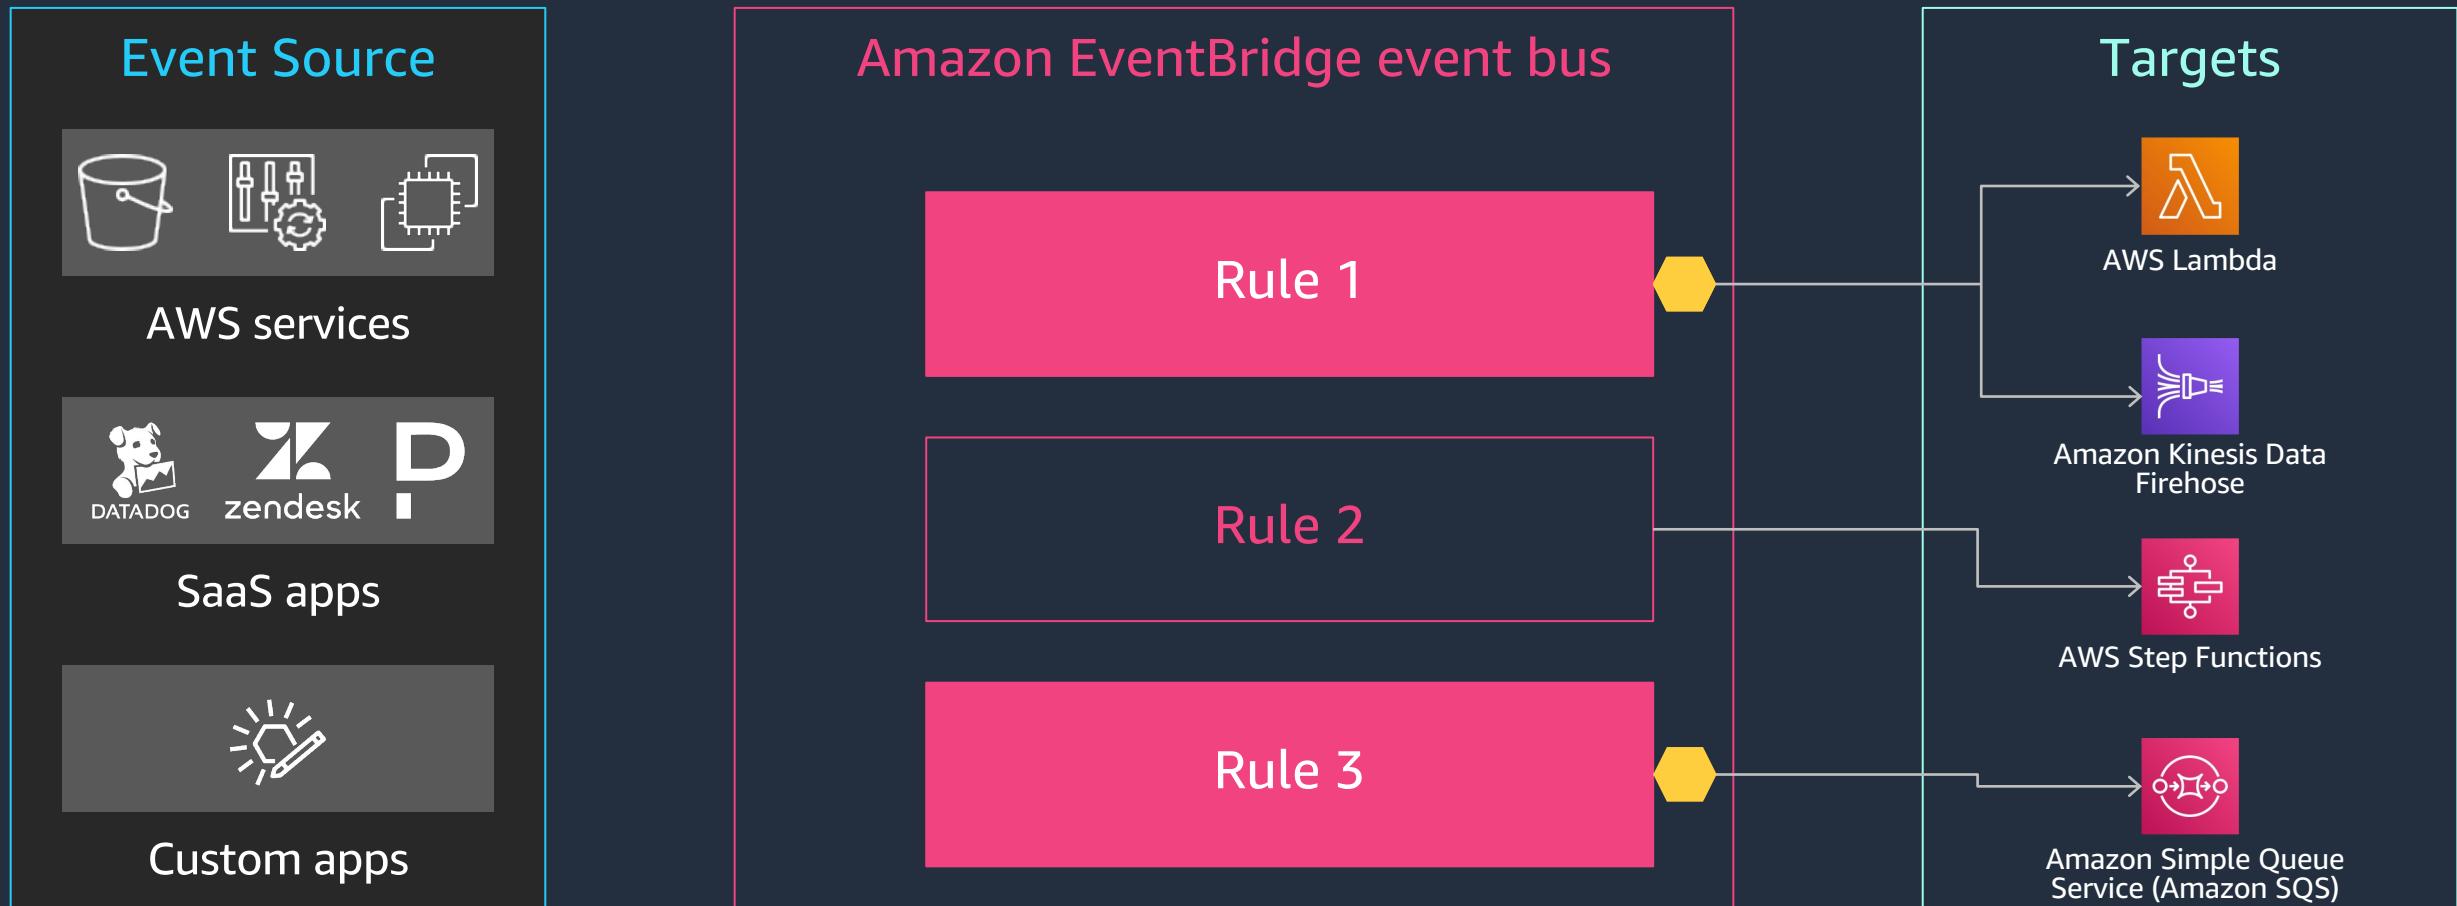
EventBridge architecture



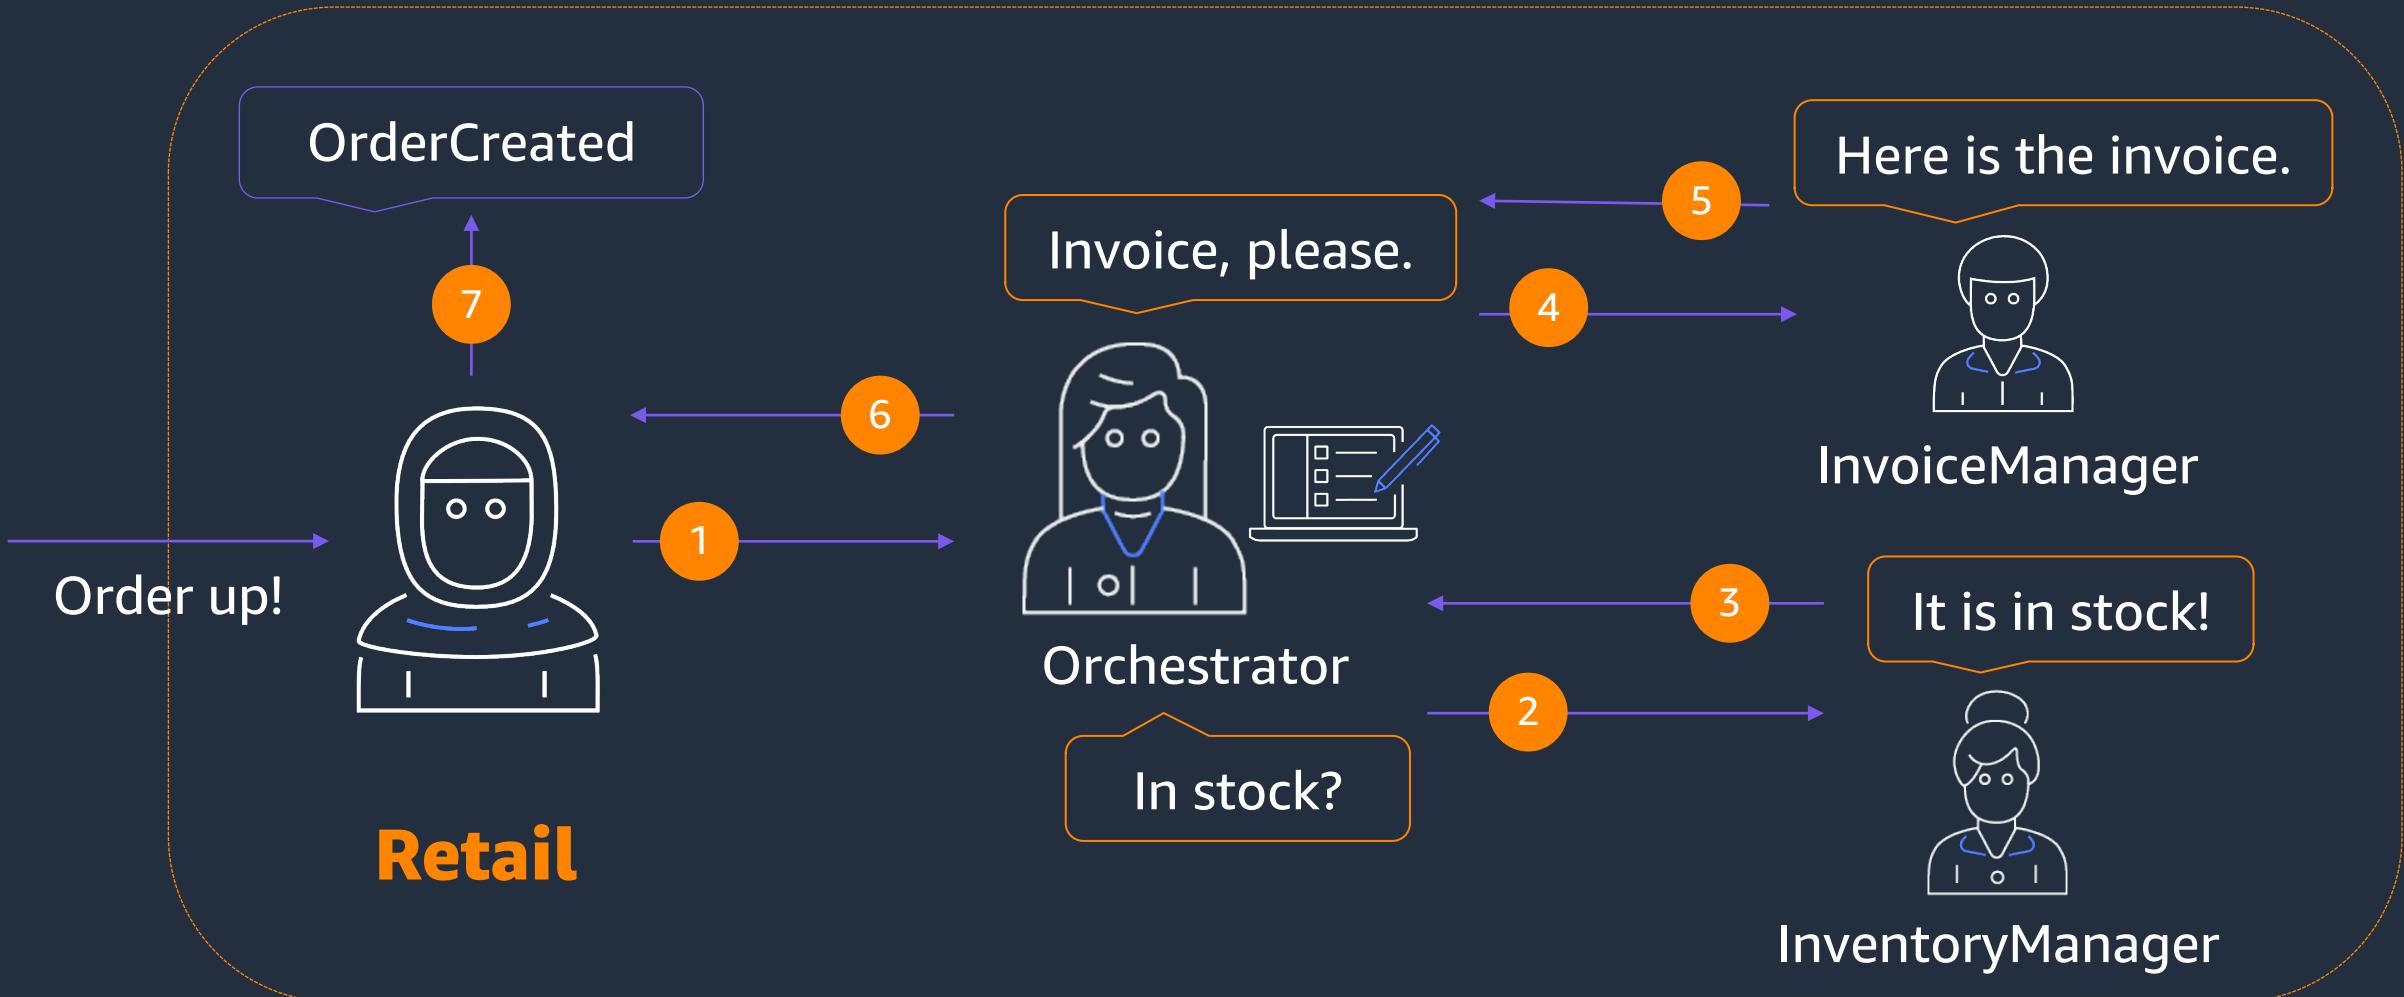
Architecture summary



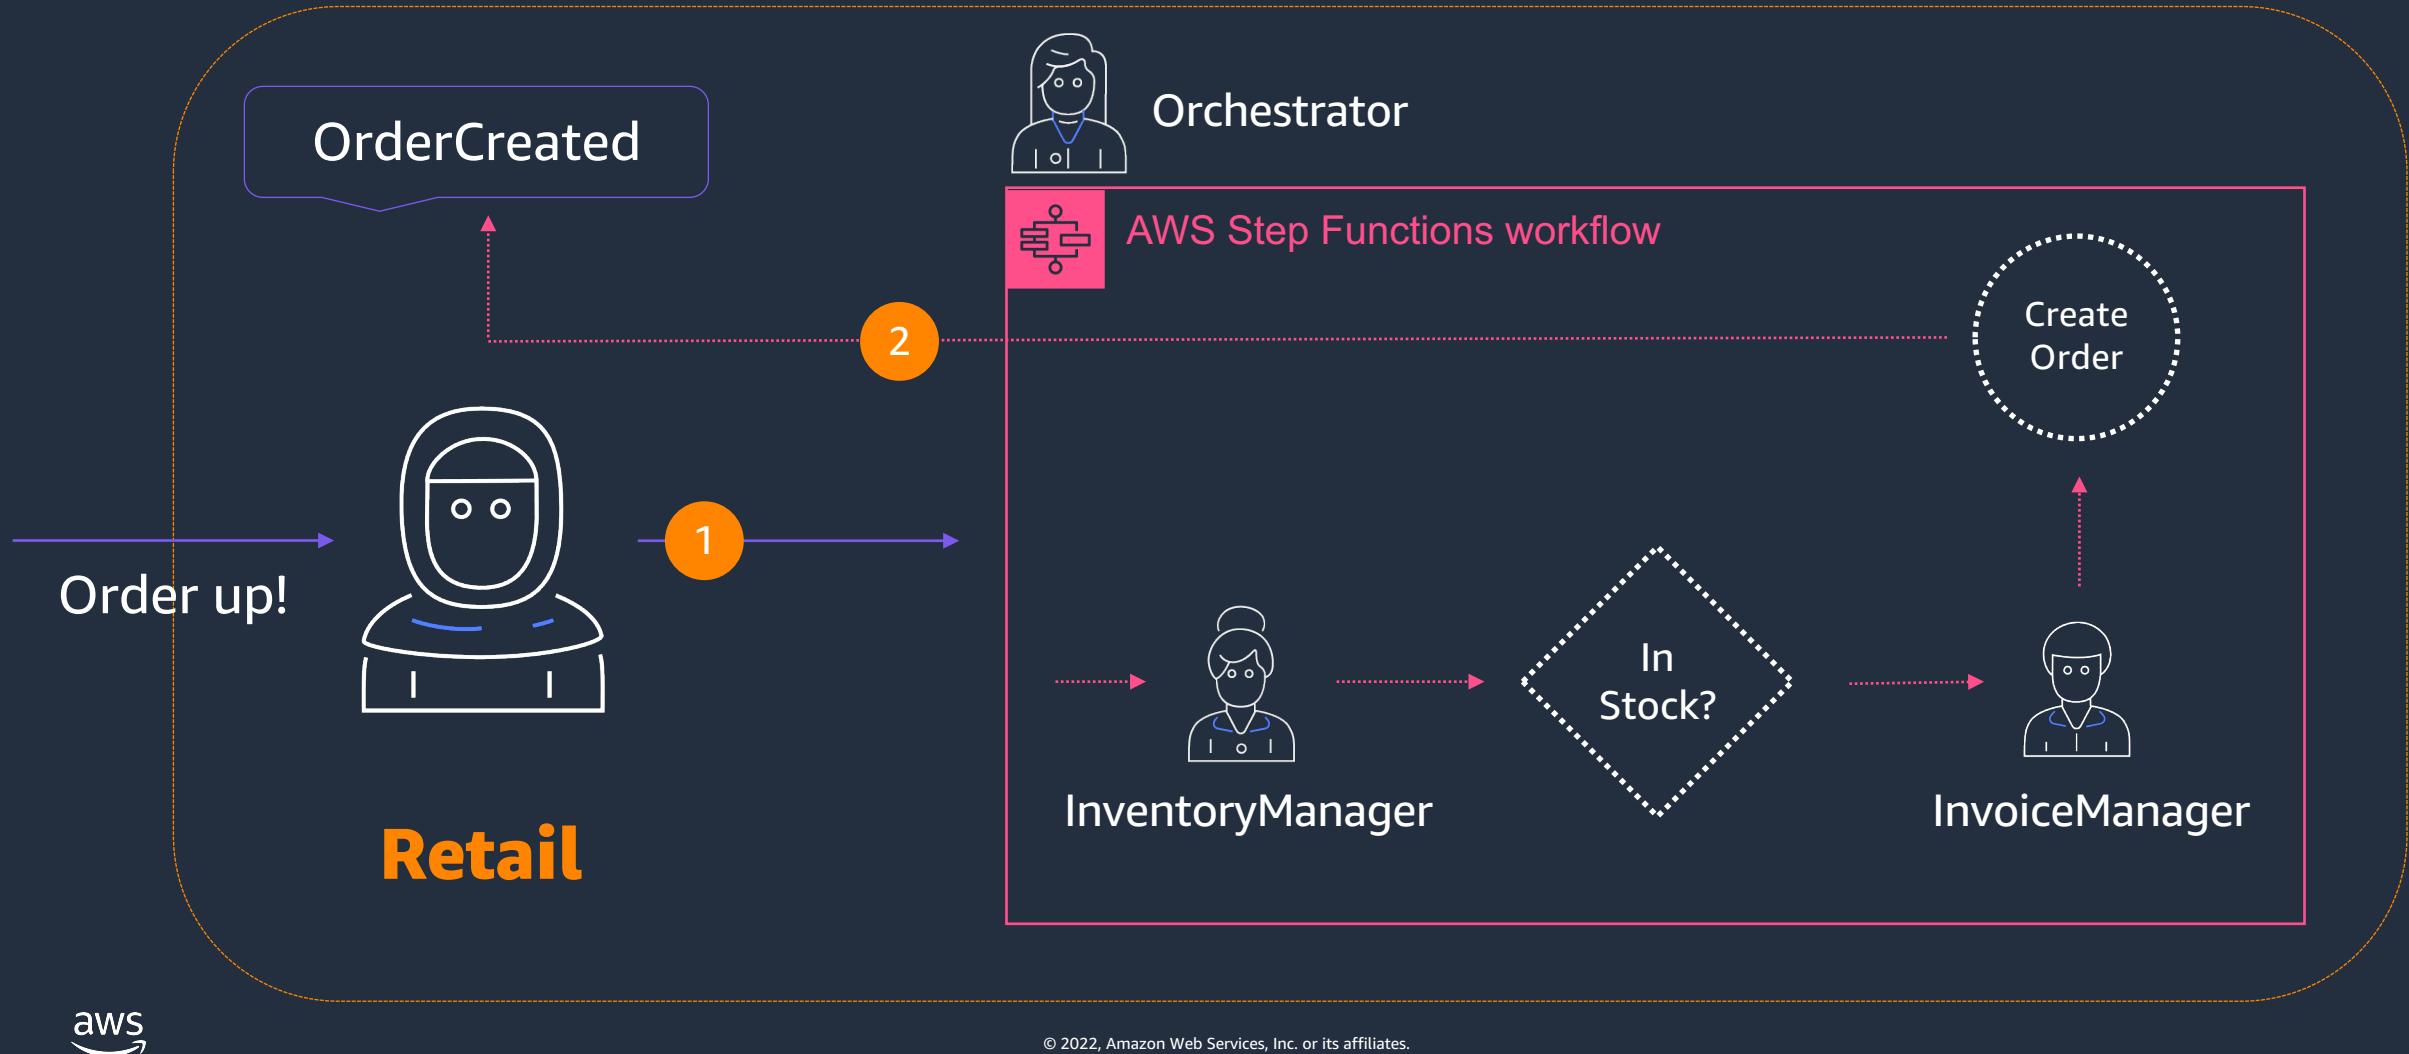
Architecture summary



Orchestrate a business process *within a domain*, resulting in a published event



Orchestrate a business process *within a domain*, resulting in a published event





AWS Step Functions



The **workflows** you build with Step Functions are called **state machines**, and **each step** of your workflow is called a **state**.



When you execute your state machine, each move from one state to the next is called a **state transition**.



You can **reuse components**, easily edit the sequence of steps or swap out the code called by task states as your needs change.

Step Functions workflow designer (Preview)

Configuration Input Output Error handling

Retry on errors - optional Info

Catch errors - optional Info

Timeout - optional Info

Heartbeat - optional Info

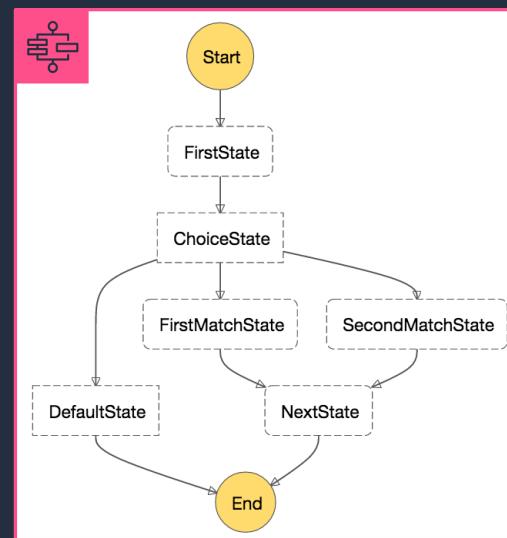


Visual workflows

Define

- Drag-and-drop with Workflow Studio
- JSON (Amazon States Language)
- CDK (TypeScript, JavaScript, Python, Java, C#)
- Data Science SDK (Python)

Visualize



Execute and monitor

Execution Arm: arn:aws:states:eu-central-1:492419596455:execution:Orderer>New_Order
New_Order ✓

Graph | Code

Execution Status: Succeeded

State Machine ARN: arn:aws:states:eu-central-1:492419596455:stateMachine:Orderer

Execution ID: arn:aws:states:eu-central-1:492419596455:execution:Orderer_New_Order

Started: Nov 20, 2016 9:58:28 AM

Closed: Nov 20, 2016 9:58:32 AM

ID Type

- 1 ExecutionStarted
- 2 TaskStateEntered
- 3 LambdaFunctionScheduled

CloudWatch > Log Groups > /aws/states/Tmp212262019-Logs > states/Tmp212262019/2019-11-26-17-30/78172538

Filter events

Time (UTC +00:00) Message

2019-11-26

17:32:13 {"id": "1", "type": "ExecutionStarted", "details": {"roleArn": "arn:aws:iam::██████████:role/service-role/Temp211262019", "input": "\n\n"}, "previousEventId": "", "eventTimestamp": "1574769533425", "executionArn": "arn:aws:states:ca-central-1:██████████:execution:Temp212262019:9f0Fa2e2-ada7-74b9-40be-0b468bf057b9:c2873d15-496f-42d2-acd7-4"}

Expand all



Step Functions integration types

Optimized integrations

Customized to simplify the usage of 17 AWS services

Supported Integration patterns:

- Request Response
- Wait for a Callback (.waitForTaskToken)
- Run a Job (.sync)

AWS SDK integrations

Call 200 AWS services directly (9000+ API actions)

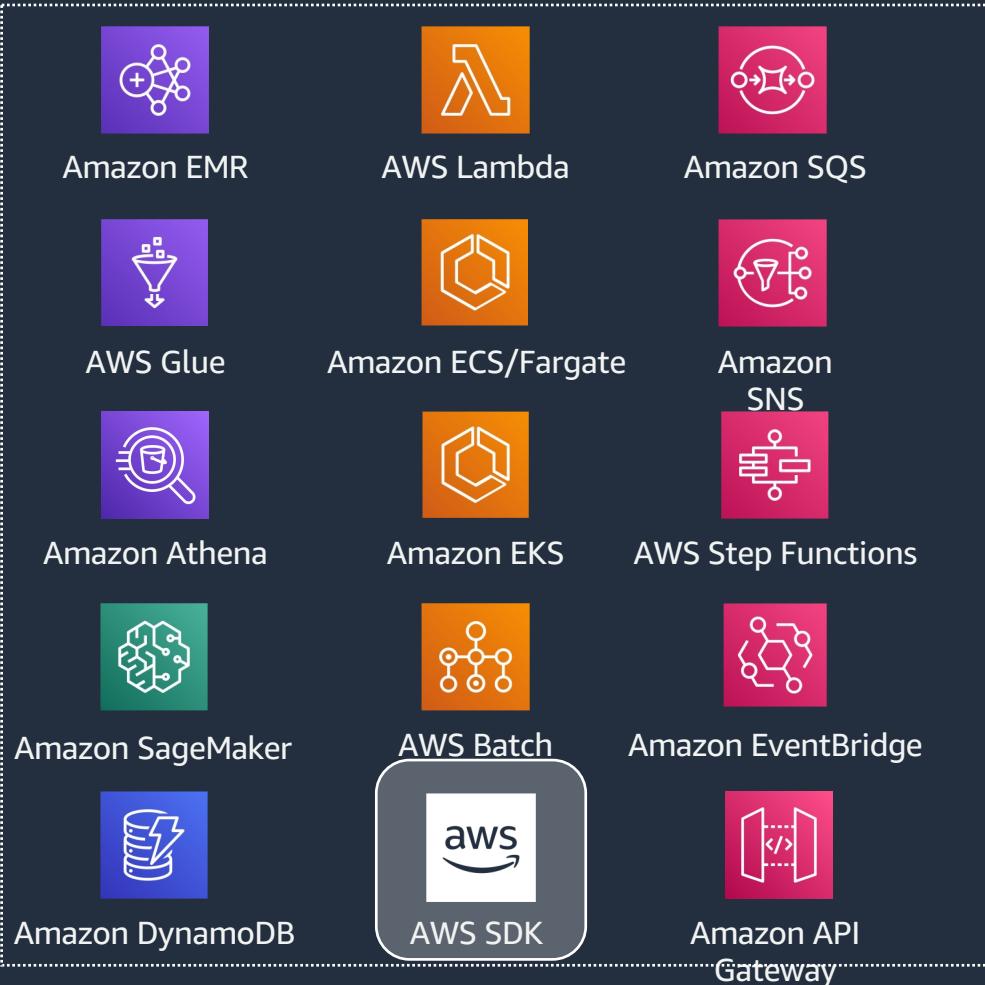
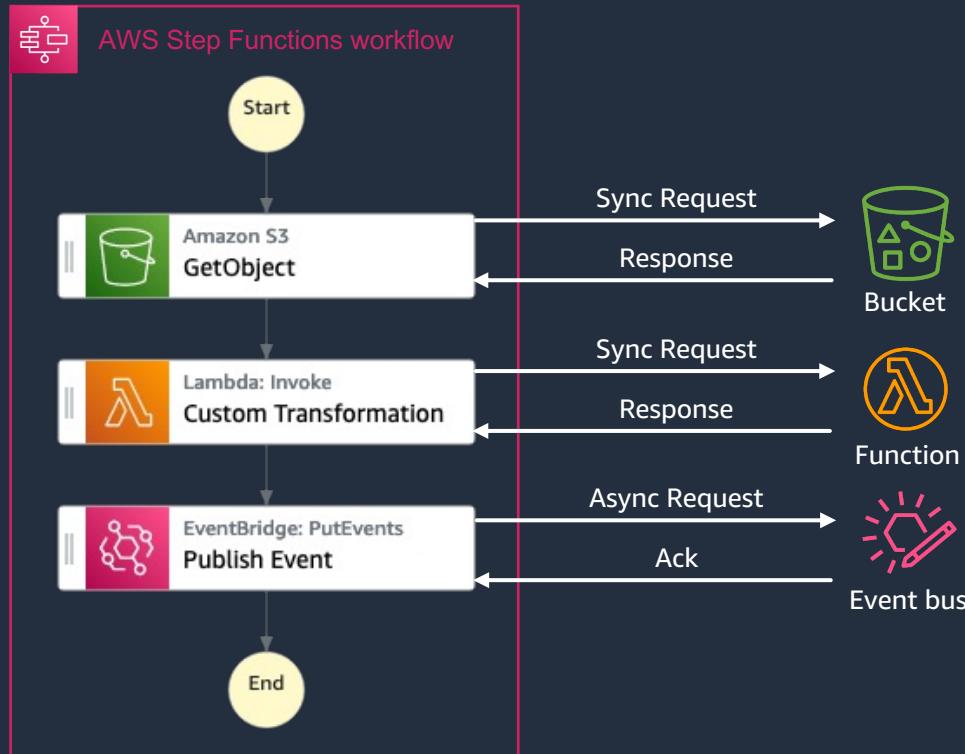
Supported Integration patterns:

- Request Response
- Wait for a Callback (.waitForTaskToken)



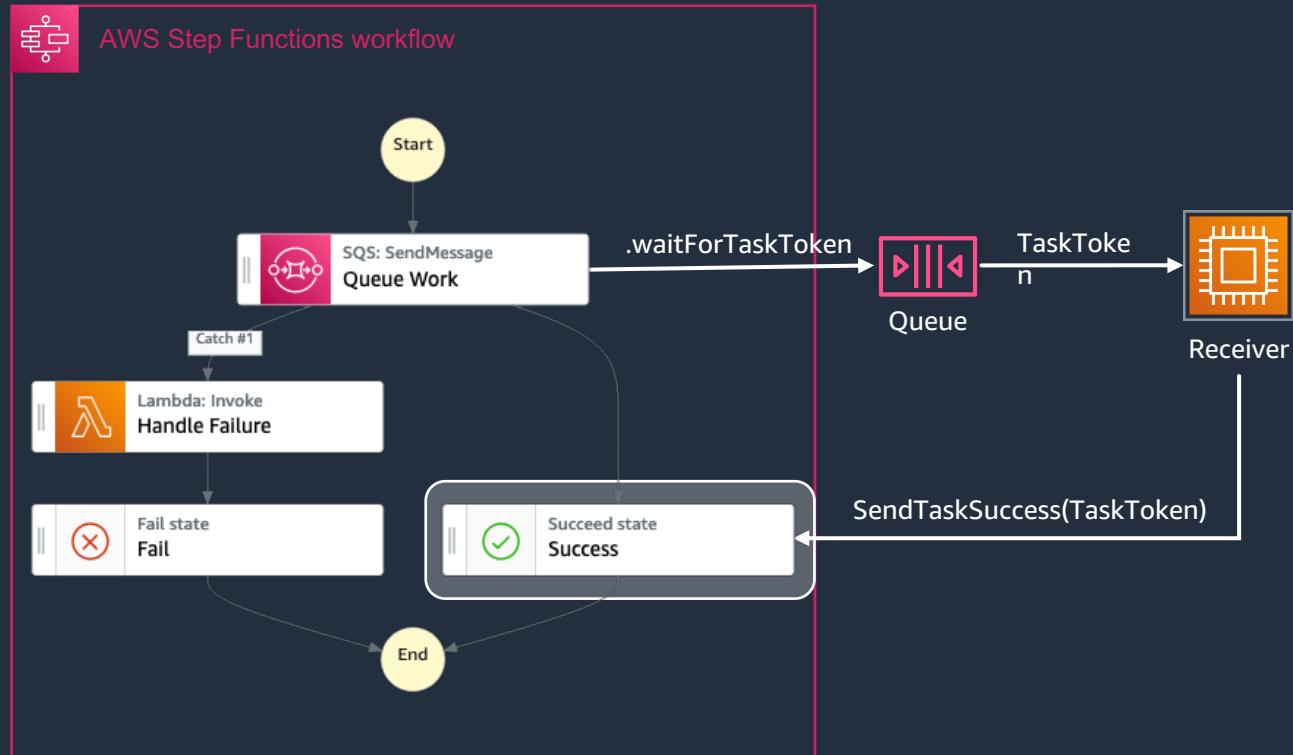
Step Functions: Request Response

Service Integration Pattern



Step Functions: Wait for a Callback (.waitForTaskToken)

Service Integration Pattern



AWS Lambda



Amazon SQS



Amazon ECS/Fargate



Amazon SNS



AWS Step Functions



Amazon EventBridge



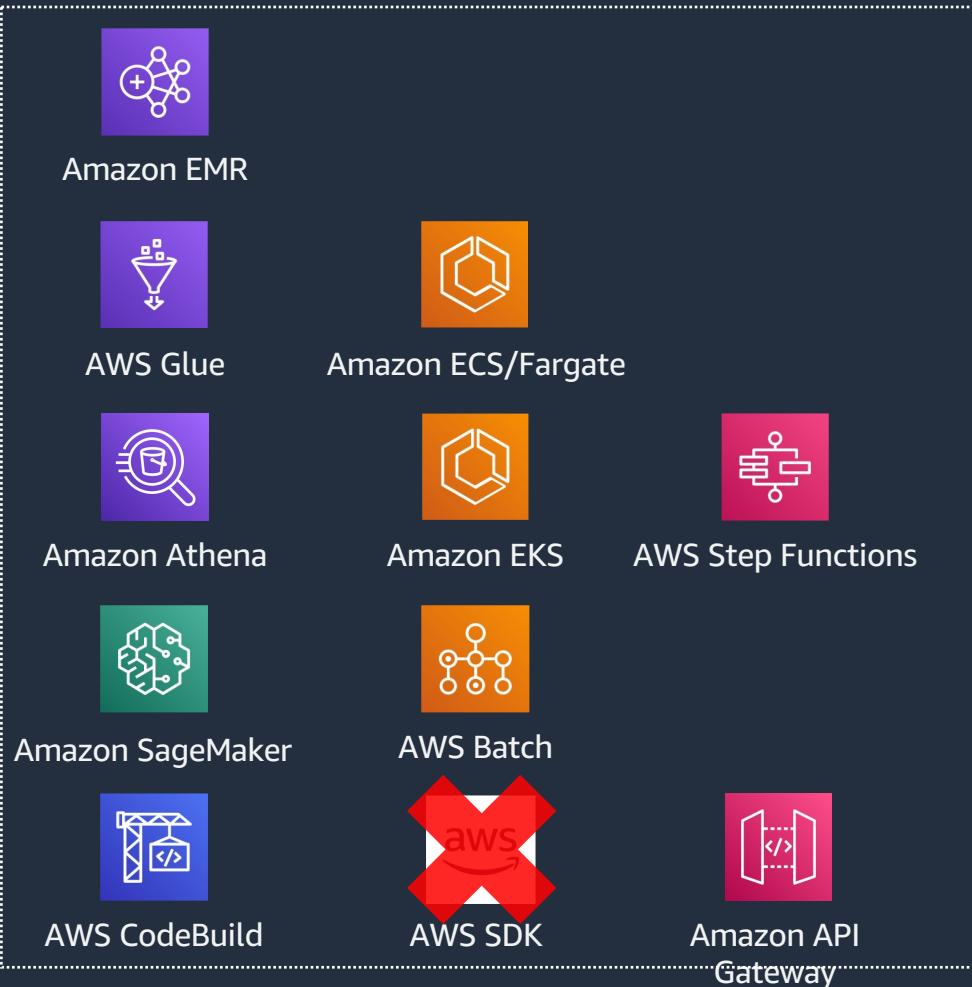
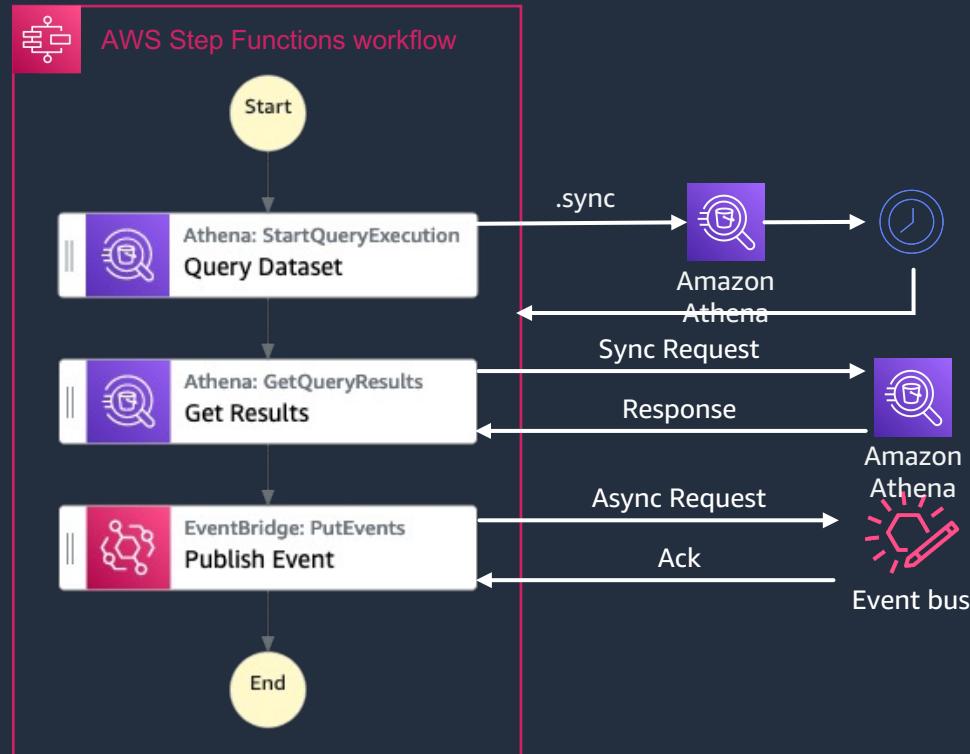
AWS SDK



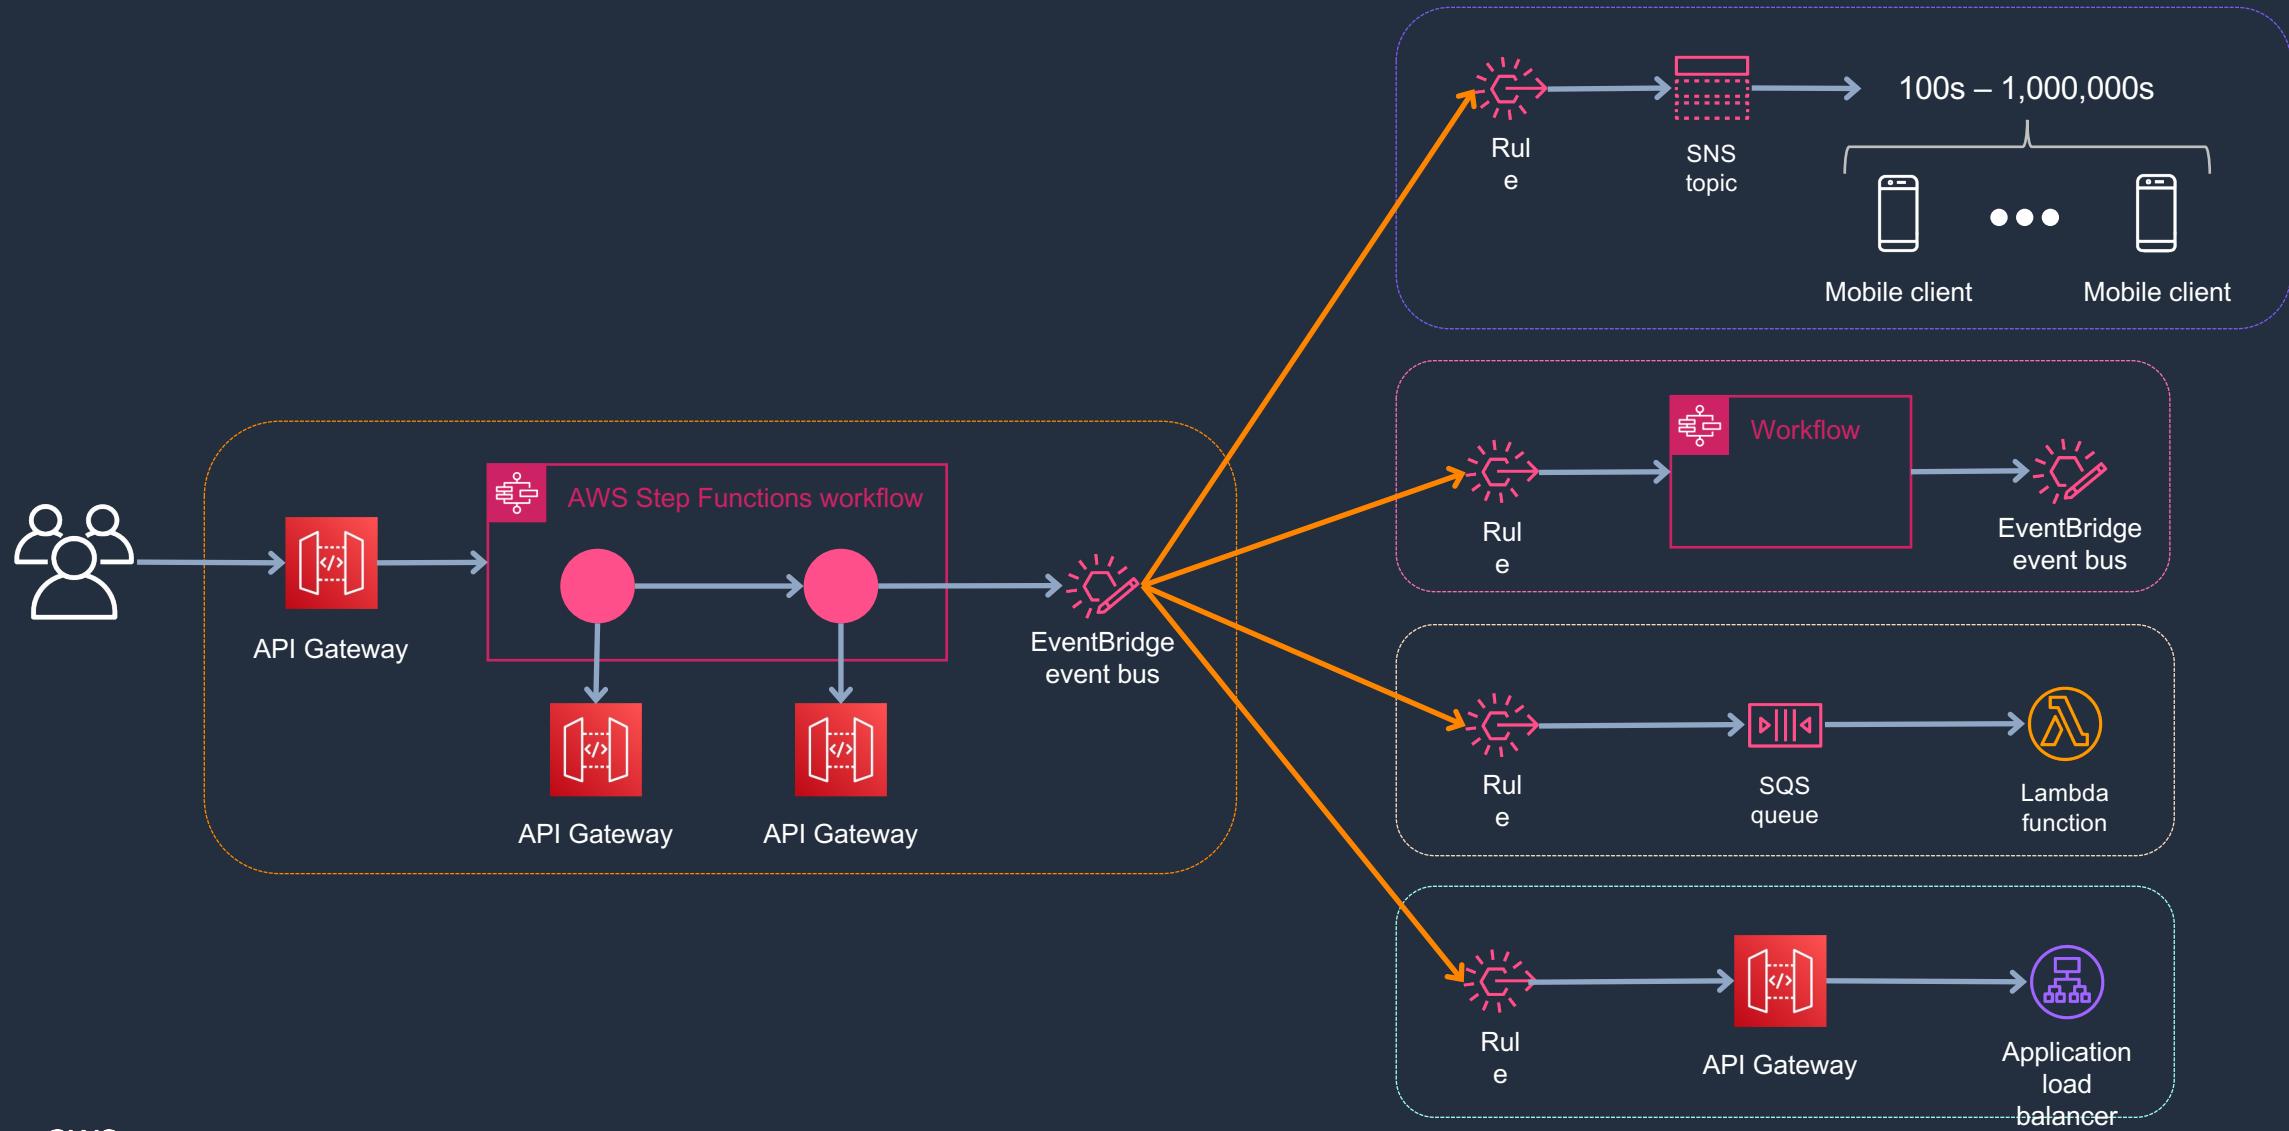
Amazon API Gateway

Step Functions: Run a Job (.sync)

Service Integration Pattern



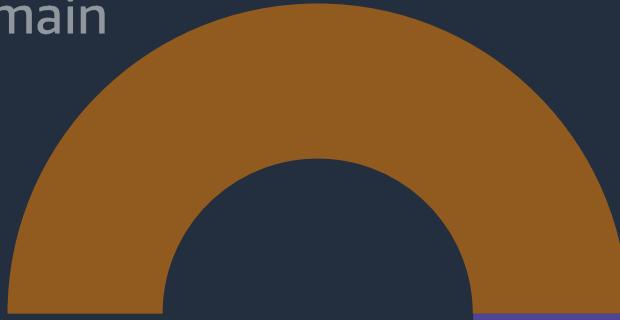
Better together: Orchestration + Choreography



Journey to event-driven architectures

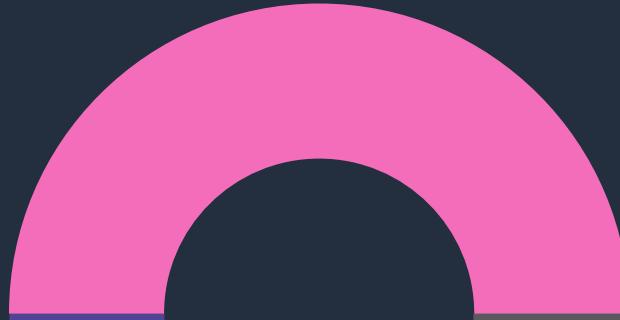
STEP 1

Start with the
domain



STEP 3

Pick an event store



STEP 2

Coordinate events



Improve resiliency and scalability with event stores



Buffer messages until service are available to process

Event stores handle messages and streams

Message Processing



- The individual **message** is the unit of work
- Computation/processing per message
- Message occurrence varies
- Message are deleted after consumption
- **No need to track the position**
- **DLQ functionality built-in**

Stream Processing



- The message **stream** is the unit of work
- Complex computation on many messages
- Constant stream of messages
- Messages are available after consumption until expiration
- Each **client** needs to track the current position in the stream
- **No built-in DLQ functionality**



Amazon Simple Queue Service (SQS)

WHAT IT IS

Simple, flexible, fully managed message queuing service for reliably and continuously exchanging any volume of messages from anywhere

(Standard and FIFO)

USE CASE

Build decoupled, highly scalable microservices, distributed systems, and serverless applications in the cloud

COOL CAPABILITIES

Nearly infinite scalability without pre-provisioning capacity



Event Source



25B messages per hour





Amazon Kinesis

WHAT IT IS

Enables you to **ingest**, **buffer**, and **process** streaming data in **real-time**, so you can derive **insights in seconds** or minutes instead of hours or days.

USE CASE

Ingest real-time data such as video, audio, application logs, website clickstreams, and IoT telemetry data for machine learning, analytics.

COOL CAPABILITIES

Can handle any amount of streaming data and process data from hundreds of thousands of sources with very low latencies.



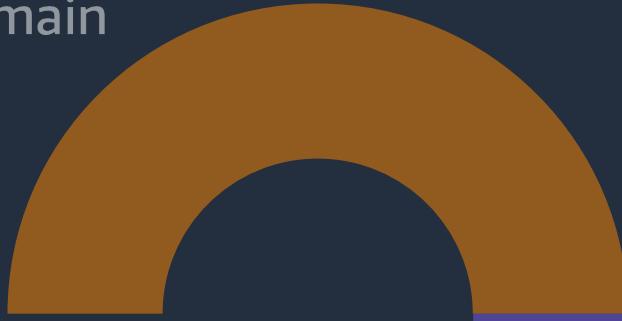
Event Source



Journey to event-driven architectures

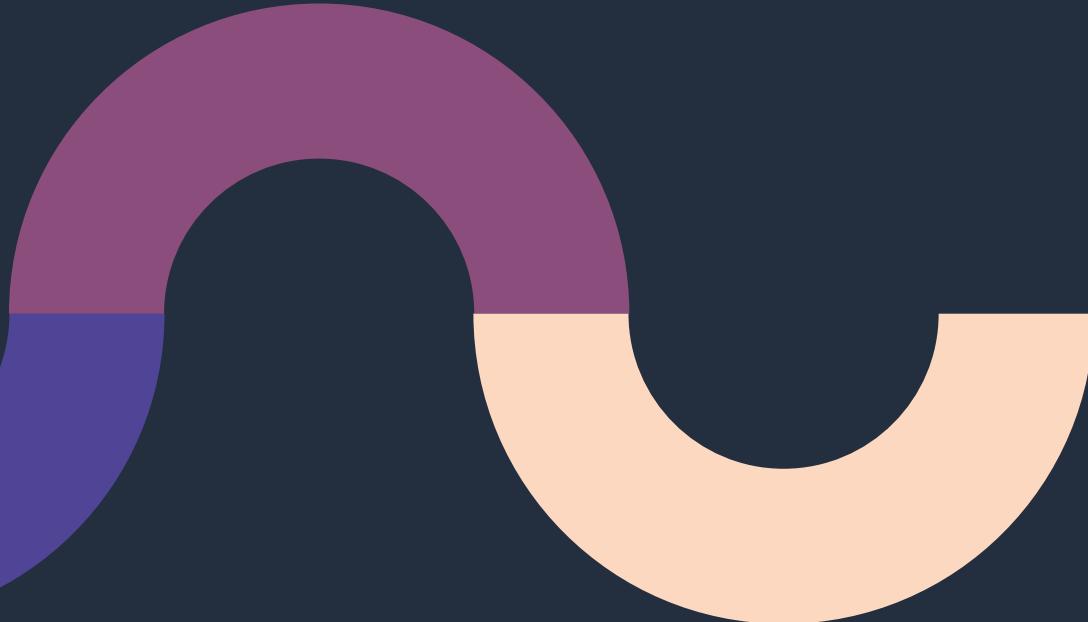
STEP 1

Start with the
domain



STEP 3

Pick an event store



STEP 2

Coordinate events

STEP 4

Structure events

Sparse events vs. full state descriptions

Order 123 was created at 10:47 a.m. by customer 456



Events

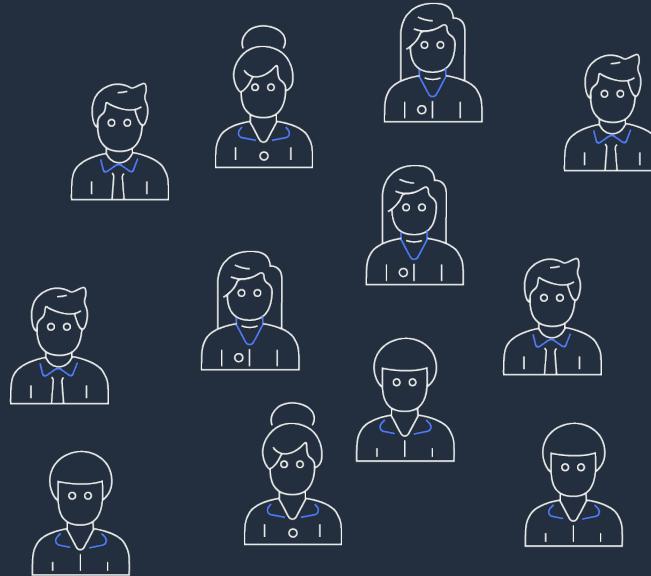
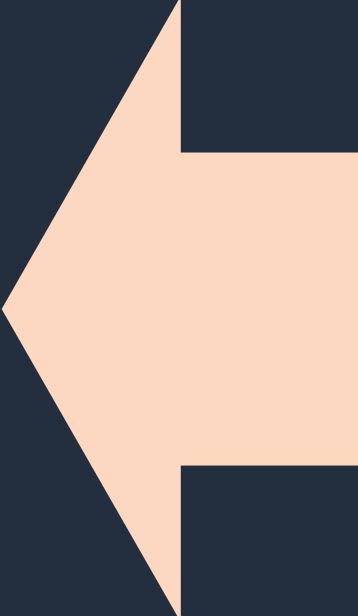
Order 123 was created at 10:47 a.m. by customer 456. The current status is Open, the total was \$237.51, the items were ...



Full state description

Considerations with sparse events

Order 123 was
created by
customer 456



What are the
details for
order 123?

Considerations with full state descriptions

```
{  
  "source": "com.orders",  
  "detail-type": "OrderCreated",  
  "detail": {  
    "metadata": {  
      "idempotency-key": "c1b95b88"  
    },  
    "data": {  
      "order-id": "1073459984",  
      "status": "Open",  
      "total": "237.51"  
    }  
  }  
}
```

- Event schemas should be backwards compatible

Considerations with full state descriptions

```
{  
  "source": "com.orders",  
  "detail-type": "OrderCreated",  
  "detail": {  
    "metadata": {  
      "idempotency-key": "c1b95b88"  
    },  
    "data": {  
      "order-id": "1073459984",  
      "status": "Open",  
      "total": "237.51"  
    }  
  }  
}
```

- Event schemas should be backwards compatible
- Cost to calculate values can increase over time



Considerations with full state descriptions

```
{  
  "source": "com.orders",  
  "detail-type": "OrderCreated",  
  "detail": {  
    "metadata": {  
      "idempotency-key": "c1b95b88"  
    },  
    "data": {  
      "order-id": "1073459984",  
      "status": "Open",  
      "total": "237.51"  
    }  
  }  
}
```

- Event schemas should be backwards compatible
- Cost to calculate values can increase over time



Design Considerations



Delivery semantics

At-least once delivery

Events can be delivered to a target more than once. Include logic to detect duplicate events by tracking the state of processed events (**idempotent**).

- Amazon EventBridge
- Amazon SNS Standard
- Amazon SQS Standard
- Amazon Kinesis

Exactly-once delivery*

Specify an identifier used by the AWS service for **deduplication**.

- Amazon MQ
- Amazon MSK
- Amazon SNS FIFO
- Amazon SQS FIFO

** outside of application error handling*



Event uniqueness

Idempotency

Operation will return the same results whether it is called once or multiple times (EIP).

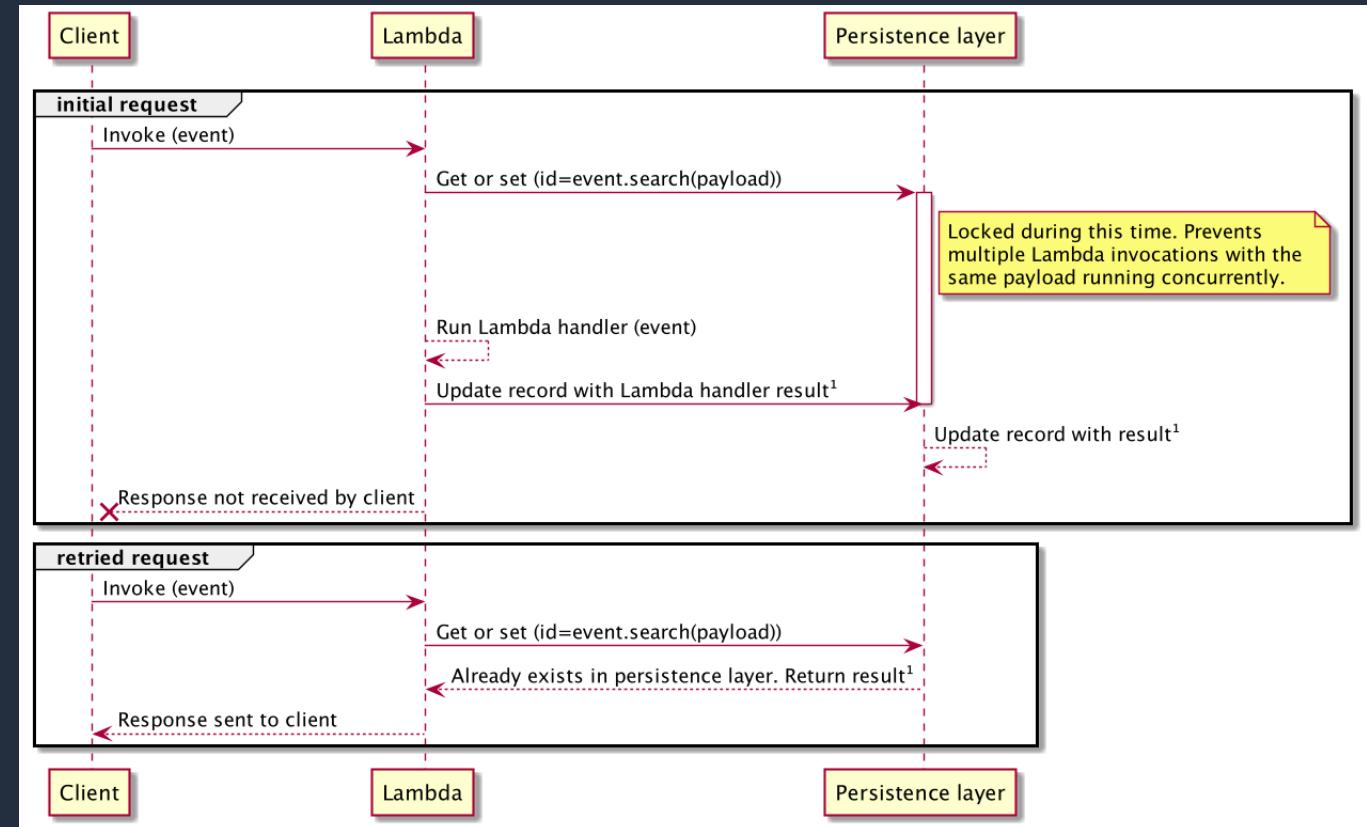
Idempotency Key

Assigned to the message by the sender to simplify deduplication by the receiver.

```
{  
  "source": "com.orders",  
  "detail-type": "OrderCreated",  
  "detail": {  
    "metadata": {  
      "idempotency-key": "c1b95b88"  
    },  
    "data": {  
      "order-id": "1073459984"  
    }  
}
```

Event deduplication strategies

- Manage idempotent state in an external data store
- Specify a time window for deduplication
- Lambda Powertools Python Idempotency library
- Making retries safe with idempotent APIs (AWS Builders' Library)



(*Lambda Powertools Python Idempotency*, <https://bit.ly/3iWKu7z>)

Ordering semantics*

Unordered

Events can be delivered out of order.

- Amazon EventBridge
- Amazon SNS Standard
- Amazon SQS Standard

Ordered

Events are delivered in order within a partition, message group, etc (no global order).

- Amazon MQ
- Amazon MSK
- Amazon Kinesis
- Amazon SNS FIFO
- Amazon SQS FIFO

** Out-of-order event handling logic is highly application specific. If the application cannot be designed to handle out-of-order events, consider orchestration instead.*



Next Steps



- Event-driven architecture best practices
<https://aws.amazon.com/event-driven-architecture/>
- Digital Course: Architecting Serverless Solutions
<https://www.aws.training/Details/eLearning?id=42594>
- Event-driven reference architecture
<https://github.com/aws-samples/aws-serverless-commerce-platform>
- Building event-driven architectures on AWS workshop
<http://event-driven-architecture.workshop.aws>





Thank you!

Peterson Larentis

linkedin.com/in/peterson-larentis/

Repo Git with Deck + Additional Materials



<https://github.com/plarentis/tdc-2022>

🙏 Survey 🙏



<https://eventbox.dev/survey/S8YLY3F>